



Bayesian Linear Models

The **blim** Package

```
> library(devtools)
> install_github("vankesteren/blim")
> library(blim)
```

1. The blim Framework

Anyone who can use the `lm` function in R can use `blim`! This is how to quickly make a **blimfit** object and see the results stored inside this object:

```
> fit <- blim(formula = y~x1+x2,
              data = dataset)
> summary(fit)
```

A `blimfit` object contains:

1. The **trace**: a matrix of all the sampled parameters
2. The **summary**: parameter estimates, standard errors, credible intervals, R squared and markov chain error
3. The **priors**: the prior information about the parameters
4. The **data**: matrix representations of the input data
5. The **autocorrelation**: to lag 40 for easy plotting

2. Wrapper Function

The `blim` function is a wrapper function that performs the following tasks:

1. Translate the model in formula form to model matrices readable for the samplers (see 3)
2. Optionally calculate a minimum training sample for an ANOVA model (option `mtsprior = TRUE`)
3. Parse priors input in the form `"dnorm(0,1000)"` for processing in the samplers
4. Call the sampler specified in the method-argument on the model matrices to perform analysis (see 3)
5. Process the trace returned from the sampler:
 - a. Apply burnin
 - b. Apply thinning
 - c. Create a summary with mean, se, CI and mc error
 - d. Give column names to the trace
6. Return a `blimfit` object with the indicated elements.

3. The Samplers

`blim`'s core is built around **4 samplers**:

1. CPPBR

The standard method is "cppbr": Bayes Regression in C++. This method is near-instant because of its use of the Armadillo linear algebra library in the C++ language via the Rcpp package. Bayes Regression samples the beta parameters at once from a multivariate normal distribution:

$$\beta_n \sim \mathcal{N}(\mu_n, \Lambda_n), \text{ where}$$
$$\Lambda_n = (X^T X + \Lambda_0), \text{ and}$$
$$\mu_n = (\Lambda_n)^{-1} (X^T X \hat{\beta} + \Lambda_0 \mu_0)$$

```
// Sample from multivariate normal
arma::mat Y = arma::randn(1, k);
arma::vec betan = (posterior_mu.t() +
  Y * arma::chol(posterior_sigma)).t();
```

2. RBR

The "rbr" method is an R implementation of the "cppbr" method and is included for compatibility.

3. RGS

Method "rgs" is the Gibbs sampler of `blim`. It can handle any amount (k) of predictors (v) due to the following for-loop within each iteration:

```
for (v in 1:k){
  p_tau <- 1/(t(X[,v])%*%X[,v]/var[i+1] +
    (prior_tau[v]^2))
  p_mu <- (prior_mu[v]*(prior_tau[v]^2) +
    (y-bhat[-v]%*%t(X[, -v]))%*%X[,v]/var[i+1])
    * p_tau
  bhat[v] <- rnorm(1, mean = p_mu, sd = sqrt(p_tau))
}
```

4. RMHS

The Metropolis-Hastings procedure is explained in 4.

5. The Example

Birthweight data. 189 babies were weighed. Their mothers were checked for **smoking** during pregnancy and history of **hypertension**.

		No Smoking	Smoking
	HY	μ_1	μ_2
	HY	μ_3	μ_4

In addition, **mother's weight** at the last menstrual period was measured. This could act as a covariate. Using the table above, I can specify an **informative hypothesis** about the data: I expect $\mu_1 > \mu_2 > \mu_3 > \mu_4$, based on my "extensive" medical background. To test the hypothesis, I specify an **ANOVA**-like model with dummy variables:

$$bwt_i = \mu_1 \cdot no_i + \mu_2 \cdot sm_i + \mu_3 \cdot hy_i + \mu_4 \cdot smhy_i + \varepsilon_i$$

In the `blim` package, this simply translates to:

```
mod1 <- blim(bwt~0+no+sm+hy+smhy,
             data = birthwt,
             iter = 99999,
             mtsprior = TRUE,
             method = "rmhs")
```

4. RMHS

`blim`'s workhorse "rmhs" method uses **Gibbs sampling** just like in method "rgs" for conjugate priors and a **random-walk Metropolis** algorithm for evaluating the posterior in case of **nonconjugate priors**:

```
for (v in 1:k){
  if (fun_prior_b[v] == "dnorm"){
    # Gibbs procedure if prior is conjugate
  } else {
    # Likelihood variance and mean
    l_tau <- var[i+1]/t(X[,v])%*%X[,v]
    l_mu <- (y-bhat[-v]%*%t(X[, -v]))%*%X[,v]/
      var[i+1] * l_tau
    # posterior <- prior * likelihood
    fx <- function(x){
      eval(parse(text=prior_beta[v])) *
        dnorm(x, mean = l_mu, sd = sqrt(l_tau))
    }
    # Sample a candidate value
    bstar <- bhat[v] + rnorm(1, 0, tune[v])
    # Compute the acceptance ratio
    r <- min(1, fx(bstar)/fx(beta[i,v]))
    # Assign when accepted
    if (runif(1) <= r) bhat[v] <- bstar
  }
}
```

8. Posterior Predictive Check

An important assumption of the ANOVA is the **normality of residuals**. `blim` can construct a p-value for this using the kolmogorov-smirnov statistic as its **discrepancy measure**. In R the observed discrepancy vector is calculated in this way:

```
ObsD <- apply(trace,1,function(b){
  resid <- blimfit$y-blmfit$X%*%b
  D <- ks.test(resid,pnorm)$statistic
  return(D)
})
```

For the repD, yrep is calculated as follows for each iter:

```
y <- apply(blimfit$X%*%b,1,function(mean){
  rnorm(1,mean,sqrt(var))
})
```

Calling `PPC(mod2, type = "norm")` yields a posterior predictive p-value of around 0.5, indicating no violation of the residual normality assumption.

6. Model Selection

I specify a second model with mother's weight (`lwt`) as a standardised covariate. I give it a **cauchy prior** with location 0 and scale 100, as the effect will be in terms of grams of birthweight. This seems a reasonable prior, as I don't know whether the effect is positive or negative, but I do have expectations about the approximate range.


```
mod2 <- blim(bwt~0+no+sm+hy+smhy+scale(lwt),
             data = birthwt,
             # Mu priors the same as mtsprior for mod1
             prior_b = c(rep("dnorm(2950,402403)",4),
                           "dcauchy(0,100)"),
             method = "rmhs",
             # Dynamic tuning every 100th iter (see 10)
             dtuning = 100)
```

Unlike frequentist statistics, where model selection too often happens through p-values, Bayesian analysis allows for model selection with **Bayes Factors**, quantifying the amount of evidence there is for model 1 relative to model 2. `blim` can calculate a Bayes Factor comparing two models using the **Laplace-Metropolis estimator**:

$$f(D) \approx (2\pi)^{P/2} |\mathbf{H}^*|^{1/2} f(\theta^*) f(D|\theta^*)$$

Where the quantities \mathbf{H}^* and θ^* are estimated via the trace covariance matrix and the conditional medians, respectively. It can also quantify the **precision** of this estimate using a **bootstrap**. The precision increases when the amount of iterations in the sampler increases.

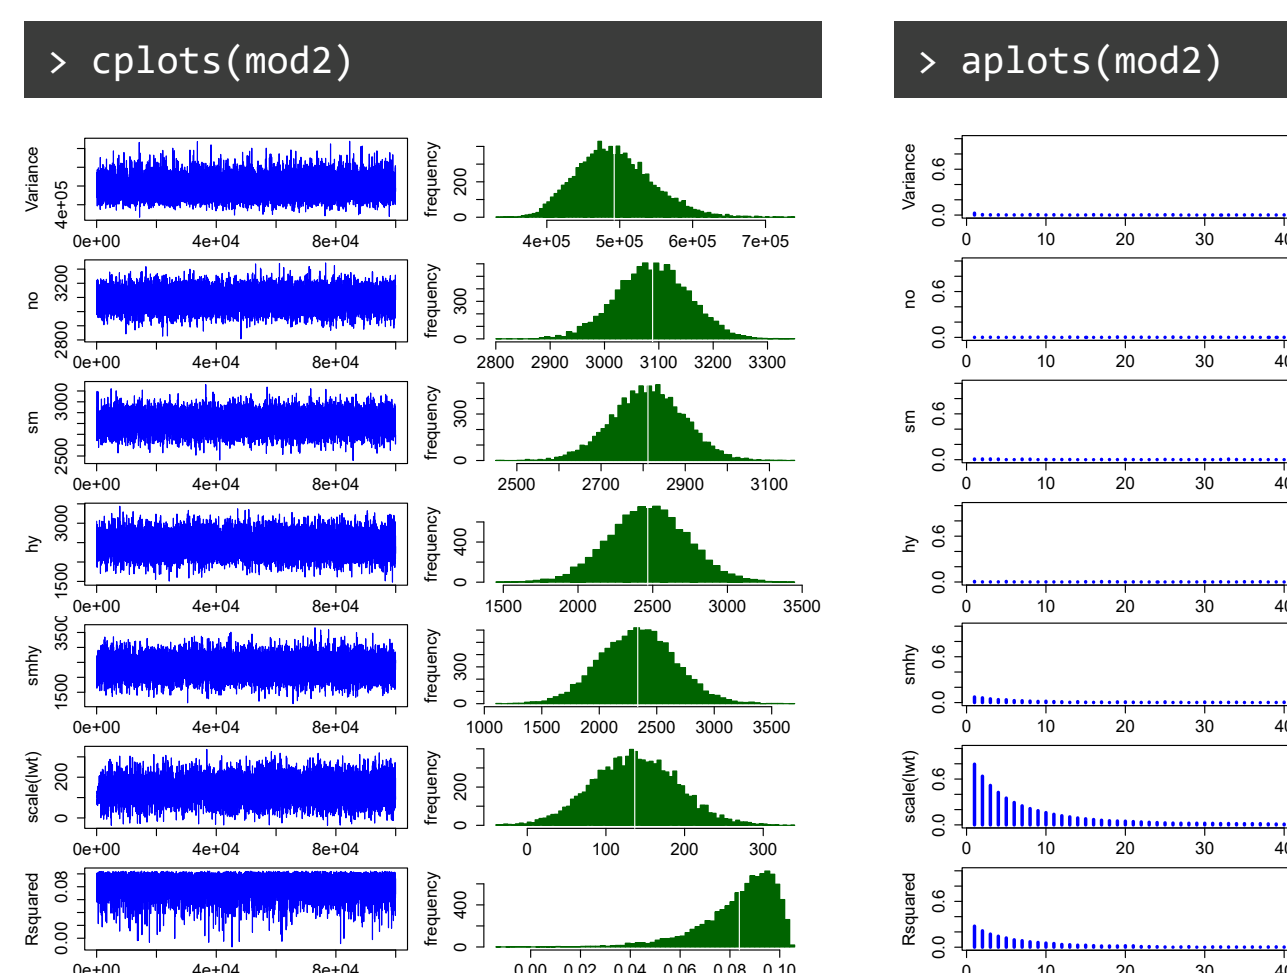
```
> BF(mod1, mod2, bootstrap = TRUE, plot = TRUE)
```



BF12 is around 0.0025. This indicates that model 2 is 400 times as likely as model 1. Here, **DIC** difference ≈ 15.7

7. Convergence

Convergence and autocorrelation for the parameters and R^2 of model 2 are shown below. Note that the trace plot indicates no lack of convergence, the autocorrelation for the μ s is low, and the autocorrelation for the covariate parameter with nonconjugate prior drops off quickly.




10. Dynamic Tuning

`blim` performs tuning of the Metropolis algorithm acceptance rate by dynamically optimising the **variance of the proposal distribution**. This variance starts at 2.38 for each conditional posterior and is regulated based on keeping the acceptance rate within a margin, where the optimal acceptance rate is set at 0.45. `blim` reports final variance.

```
# pcur is current acceptance rate, popt=0.45
if (pcur > 0.5 || pcur < 0.15){
  tune[v] <- (tune[v]*(1/pnorm(popt/2)))/
    (1/pnorm(pcur/2))
}
```

9. Conclusion



Covariate: mean 136, 95% CI [32;245]. For 1 sd increase in mother's last weight, birthwt increases by 136 grams.

Birthwt nothing (μ_1): 3089 grams, 95% CI [2957; 3221]
Birthwt smoke (μ_2): 2811 grams, 95% CI [2645; 2978]
Birthwt hypert (μ_3): 2469 grams, 95% CI [1946; 2992]
Birthwt sm+hy (μ_4): 2337 grams, 95% CI [1696; 2971]

Covariate: 136, 95% CI [32;245]. (For 1 sd increase in mother's weight, birthweight increases by 136 grams)

The means seem to have the right direction, but I want to test the hypothesis! `blim`'s inequality constraint BF estimator evaluates the hypothesis in the trace of the posterior (**fit**) and a randomly generated prior trace (**complexity**):

```
model = "par[1] > par[2] & par[2] > par[3] & par[3] > par[4]"
fit <- mean(apply(trace, 1,
  function(par) eval(parse(text=model))))
com <- mean(apply(prior, 1,
  function(par) eval(parse(text=model))))
BF <- fit/com
```

My hypothesis has a Bayes Factor of around 12, which indicates **strong evidence** relative to the unconstrained hypothesis!